

REACH-AI 2026: Community-Informed Evaluation and Benchmarking of Culturally Competent Generative AI

Jamila Smith-Loud
Google
Washington, DC, USA
jsmithloud@google.com

Lijun Qian
Prairie View A&M University
Prairie View, TX, USA
liqian@pvamu.edu

Andrew Smart
Google
Washington, DC, USA
andrewsmart@google.com

Xishuang Dong
Prairie View A&M University
Prairie View, TX, USA
xidong@pvamu.edu

Abstract

Recent advances in Generative AI (GenAI) have enabled strong capabilities in language understanding, content generation, vision-language reasoning, and multimodal interaction. However, most existing benchmarks emphasize general-purpose reasoning and task performance, offering limited evaluation of whether AI systems can accurately understand, represent, retrieve, and generate culturally grounded knowledge across diverse communities and social contexts. This limitation restricts the reliable deployment of GenAI in high-impact domains such as education, healthcare, and public services. The REACH-AI 2026 Workshop brings together researchers, industry leaders, practitioners, and community stakeholders to advance community-informed approaches for evaluating culturally competent Generative AI. Motivated by the REACH GenAI Consortium, a collaboration between Google Research and twelve Historically Black Colleges and Universities (HBCUs), the workshop focuses on culturally grounded datasets, community-informed taxonomies, evaluation frameworks, benchmark development, and human-centered assessment protocols for measuring cultural competency in AI systems. It aims to foster interdisciplinary dialogue across the CIKM, information retrieval, knowledge management, multimodal AI, and responsible AI communities. Through invited talks, paper presentations, and panel discussions, participants will explore methods for assessing cultural relevance, contextual understanding, representation, and knowledge utilization in GenAI systems. Expected outcomes include a community roadmap, collaborative research opportunities, and recommendations for future datasets, benchmarks, and evaluation standards for culturally competent AI.

Keywords

Culture/Social Bias, Generative AI, Large Language Models, REACH GenAI Consortium

Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for components of this work owned by others than the author(s) must be honored. Abstracting with credit is permitted. To copy otherwise, or republish, to post on servers or to redistribute to lists, requires prior specific permission and/or a fee. Request permissions from permissions@acm.org.

Conference acronym 'XX, Woodstock, NY

© 2018 Copyright held by the owner/author(s). Publication rights licensed to ACM.

ACM ISBN 978-1-4503-XXXX-X/2018/06

<https://doi.org/XXXXXXX.XXXXXXX>

ACM Reference Format:

Jamila Smith-Loud, Andrew Smart, Lijun Qian, and Xishuang Dong. 2018. REACH-AI 2026: Community-Informed Evaluation and Benchmarking of Culturally Competent Generative AI. In *Proceedings of Make sure to enter the correct conference title from your rights confirmation email (Conference acronym 'XX)*. ACM, New York, NY, USA, 4 pages. <https://doi.org/XXXXXXX.XXXXXXX>

1 Workshop Overview and Rationale

1.1 Overview

Generative AI systems increasingly rely on large-scale knowledge acquisition, retrieval, reasoning, and generation to support applications in education, healthcare, public services, and other socially impactful domains [1–4]. While recent advances in large language models (LLMs) and multimodal foundation models have demonstrated remarkable capabilities, existing evaluation benchmarks largely emphasize general-purpose reasoning and task performance. Limited attention has been paid to assessing whether these systems can accurately acquire, represent, retrieve, and utilize culturally grounded knowledge across diverse communities and social contexts [5].

The REACH-AI 2026 Workshop addresses this emerging challenge through a knowledge-centric perspective on culturally competent Generative AI. Motivated by the REACH GenAI Consortium, a collaborative initiative involving Google Research and twelve Historically Black Colleges and Universities (HBCUs), the workshop focuses on the development of benchmarks, evaluation frameworks, and guidelines that reflect diverse community perspectives. .

By bringing together researchers from knowledge management, information retrieval, generative AI, and responsible AI communities, the workshop aims to establish new directions for evaluating how AI systems acquire, represent, and generate culturally relevant knowledge.

1.2 Scopes

The workshop focuses on the following topics, including but not limited to:

- Information retrieval for culturally relevant content via LLMs
- Knowledge construction for underrepresented communities
- Knowledge representation for culturally sensitive contexts
- Human-AI collaboration in knowledge acquisition

- Cultural/social bias mitigation for LLMs and multimodal models
- Human-centered AI evaluation and annotation frameworks
- Explainability and transparency in AI evaluation

1.3 Relevance to CIKM

The REACH-AI workshop addresses fundamental challenges in knowledge acquisition, representation, retrieval, and evaluation that lie at the core of the CIKM community. The REACH GenAI Consortium develops community-informed taxonomies, culturally grounded datasets, benchmark suites, rater guidelines, and evaluation frameworks that enable systematic assessment of how generative AI systems acquire, represent, retrieve, and utilize cultural knowledge.

The workshop therefore provides a unique venue for connecting advances in knowledge management, information retrieval, generative AI, multimodal learning, and AI evaluation with emerging challenges in culturally competent Generative AI. By integrating community-informed methodologies with knowledge-centric AI evaluation, the workshop broadens the scope of CIKM while attracting researchers from responsible AI, human-centered AI, and evaluation communities.

2 Organizers' Information

The workshop will be jointly organized by faculty members participating in the REACH GenAI Consortium and collaborators from Google Research.

2.1 Academic Organizer

Dr. Xishuang Dong (Main Point of Contact), Prairie View A&M University

- Role: Workshop Chair and Primary Contact
- Expertise: LLMs, information retrieval, Natural Language Processing
- Responsibilities:
 - Overall workshop coordination
 - Program development
 - Program committee management
 - Communication with CIKM organizers
- Attendance: In-person

Dr. Lijun Qian, Prairie View A&M University

- Expertise: Trustworthy AI, knowledge mining, big data analytics
- Responsibilities:
 - Paper review coordination
 - Session moderation
 - Proceedings preparation
- Attendance: In-person

2.2 Industry Organizer

Jamila Smith-Loud, Google

- Expertise: Responsible AI, community-centered AI, equity assessment in generative AI
- Responsibilities:
 - Deliver an invited keynote presentation

- Bridge academia and industry

- Attendance: In-person

Andrew Smart, Google

- Expertise: Human-centered AI, AI Fairness, AI Ethics
- Responsibilities:
 - Provide technical guidance
 - Support program development
 - Promote industry-academia collaboration
- Attendance: In-person

2.3 Other Organizer

REACH Consortium Faculty Representative

- Expertise: AI ethics, cultural studies, community-engaged research
- Responsibilities:
 - Panel organization
 - Community participation initiatives
 - Dissemination activities
- Attendance: In-person

All organizers confirm availability for on-site coordination and workshop execution at CIKM 2026. Collectively, the organizing team brings expertise spanning data mining, responsible AI, educational technologies, social computing, and large-scale evaluation frameworks.

3 Workshop Organization

3.1 Duration, Format, and Expected Contributions

The proposed workshop will be a **half day event**. Its format combines traditional academic presentations with interactive community-engagement activities:

- Keynote presentations
- Research paper presentations
- Industry and community panel discussions
- Open community roadmap session

Expected contributions include:

- **Full Research Papers (8 pages):** Novel research contributions related to culturally competent AI evaluation, datasets, benchmarks, and methodologies.
- **Short Papers (4 pages):** Work-in-progress research, system demonstrations, and emerging ideas.

3.2 Expected Participation

We estimate attendance between 40 and 80 participants from:

- Academic researchers
- Industry practitioners
- Google Research participants
- HBCU faculty and students

4 Workshop Program

4.1 Tentative Schedule

Table 1 present the tentative half-day schedule for the REACH-AI 2026 workshop.

Table 1: Tentative Half-Day Schedule for the REACH-AI 2026 Workshop.

08:00–08:15	Opening Remarks and Workshop Introduction
08:15–09:00	Keynote 1: REACH GenAI Vision: Community-Driven Evaluation of Cultural Knowledge and Competency in Generative AI (Google, Tentative)
09:00–09:15	Coffee Break and Networking
09:15–10:00	Keynote 2: Responsible AI via Cultural Knowledge and Competency (Google, Tentative)
10:00–10:15	Coffee Break and Networking
10:15–11:15	Research Paper Presentation
11:15–11:45	Panel Discussion, Community Roadmap and Future Research Directions
11:45–12:00	Closing Remarks

4.2 Attendee Engagement and Special Requirements

To encourage active participation, the workshop will incorporate:

- Live audience polling via Zoom
- Interactive Q&A sessions
- Community roadmap development exercises

Additionally, the following items are required for the workshop.

- Projector and presentation system
- Wireless microphones
- Flexible seating arrangement
- Reliable Internet connectivity

5 Submission and Participation Process

All submissions either Full Papers (up to 8 pages) or Short Papers (up to 4 pages) will follow ACM formatting guidelines and be managed through the EasyChair system.

5.1 Peer Review Process

Each submission will receive at least three reviews from members of the Program Committee. Evaluation criteria include:

- Novelty
- Technical quality
- Relevance to workshop themes
- Potential impact
- Reproducibility and clarity

5.2 Program Committee

The Program Committee will be composed of researchers from:

- REACH GenAI Consortium institutions
- Google Research
- CIKM community
- Responsible AI research community

Students, industry researchers, and community-engaged scholars are especially encouraged to participate.

6 Promotion and Dissemination

The workshop will be promoted through:

- CIKM communication channels
- REACH GenAI Consortium network
- HBCUs research networks
- Google Research outreach channels

- Academic mailing lists
- Social media platforms

Workshop outcomes will be disseminated through:

- Workshop proceedings
- Public workshop website
- Open-source repositories
- Follow-up collaborative publications

A post-workshop report will summarize key findings, lessons learned, and recommendations for future research on culturally competent AI evaluation.

7 Workshop History

REACH-AI 2026 is a new workshop motivated by the emerging need for rigorous evaluation of cultural competency in Generative AI systems. It complements recent activities in the broader AI community, including workshops and networks at ACM FAccT, NeurIPS, AAAI, and CIKM that address responsible AI, fairness, evaluation, and human-centered AI.

Unlike existing workshops that primarily focus on fairness, bias mitigation, or AI safety, REACH-AI specifically emphasizes community-informed benchmark construction, cultural competency evaluation, and collaborative methodologies involving historically underrepresented communities.

The REACH GenAI Consortium provides a unique foundation for this effort through its partnership among Google Research, Google.org, and twelve HBCUs. The workshop will serve as a platform to disseminate consortium findings, establish future research directions, and foster collaborations across the information retrieval, knowledge management, machine learning, and responsible AI communities. We envision REACH-AI becoming a recurring interdisciplinary workshop series dedicated to advancing the evaluation and development of culturally competent AI systems.

Acknowledgments

This research was supported by the Google REACH program under Award No. n299661. The authors further acknowledge the use of AI-assisted tools, including ChatGPT, solely for language refinement, grammar enhancement, and spelling correction during the preparation of this manuscript.

References

- [1] Sourojit Ghosh, Pranav Narayanan Venkit, Sanjana Gautam, Shomir Wilson, and Aylin Caliskan. 2024. Do generative AI models output harm while representing

- non-Western cultures: Evidence from a community-centered approach. In *Proceedings of the AAAI/ACM Conference on AI, Ethics, and Society*, Vol. 7. 476–489.
- [2] Tiancheng Hu, Yara Kyrychenko, Steve Rathje, Nigel Collier, Sander Van Der Linden, and Jon Roozenbeek. 2025. Generative language models exhibit social identity biases. *Nature Computational Science* 5, 1 (2025), 65–75.
- [3] Giada Pistilli, Alina Leidinger, Yacine Jernite, Atoosa Kasirzadeh, Alexandra Sasha Luccioni, and Margaret Mitchell. 2024. Civics: Building a dataset for examining culturally-informed values in large language models. In *Proceedings of the AAAI/ACM Conference on AI, Ethics, and Society*, Vol. 7. 1132–1144.
- [4] Erin van Liemt, Renee Shelby, Andrew Smart, Sinchana Kumbale, Richard Zhang, Neha Dixit, Qazi Mamunur Rashid, and Jamila Smith-Loud. 2026. Cultural Perspectives and Expectations for Generative AI: A Global Survey Approach. *arXiv preprint arXiv:2603.05723* (2026).
- [5] Yuhang Wang, Yanxu Zhu, Chao Kong, Shuyu Wei, Xiaoyuan Yi, Xing Xie, and Jitao Sang. 2024. Cdeval: A benchmark for measuring the cultural dimensions of large language models. In *Proceedings of the 2nd Workshop on Cross-Cultural Considerations in NLP*. 1–16.

Received 20 February 2007; revised 12 March 2009; accepted 5 June 2009